

## Static Analysis of Bottom Tier Sensor Networks

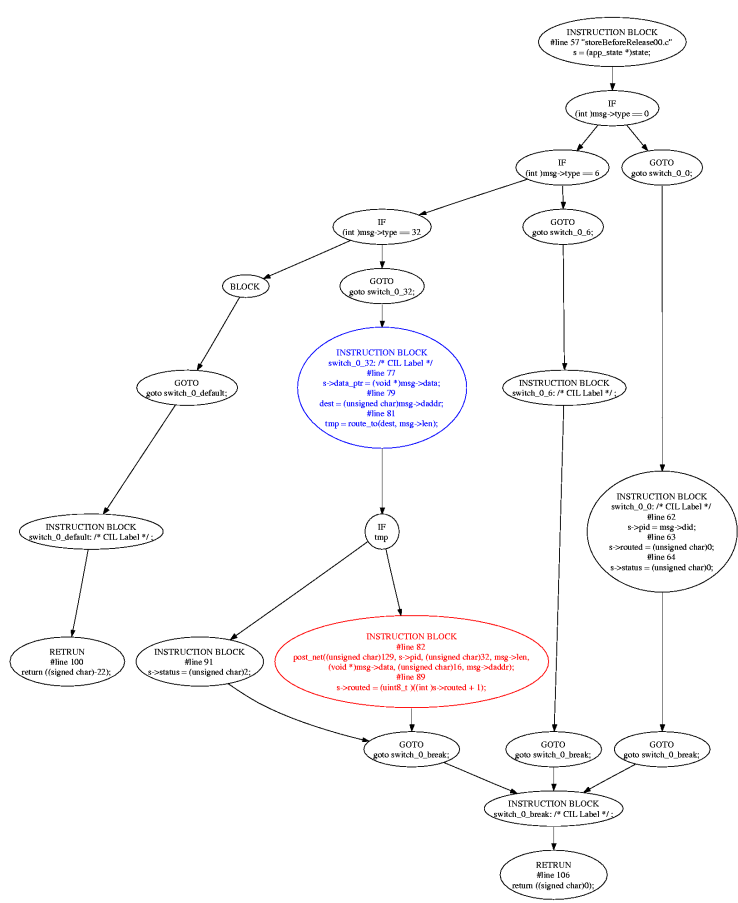
Roy Shea, Todd Millstein, Rupak Majumdar, and Mani Srivastava

Software Systems Laboratory and Networked and Embedded Systems Laboratory

### Motivation: Bringing Accessible Static Analysis To Sensor Network Development

#### Static Analysis is Cool

- Has the potential to find bugs before they arise
- Streamlines the development process
- Lowers learning curve for new developers




#### Combination of These?

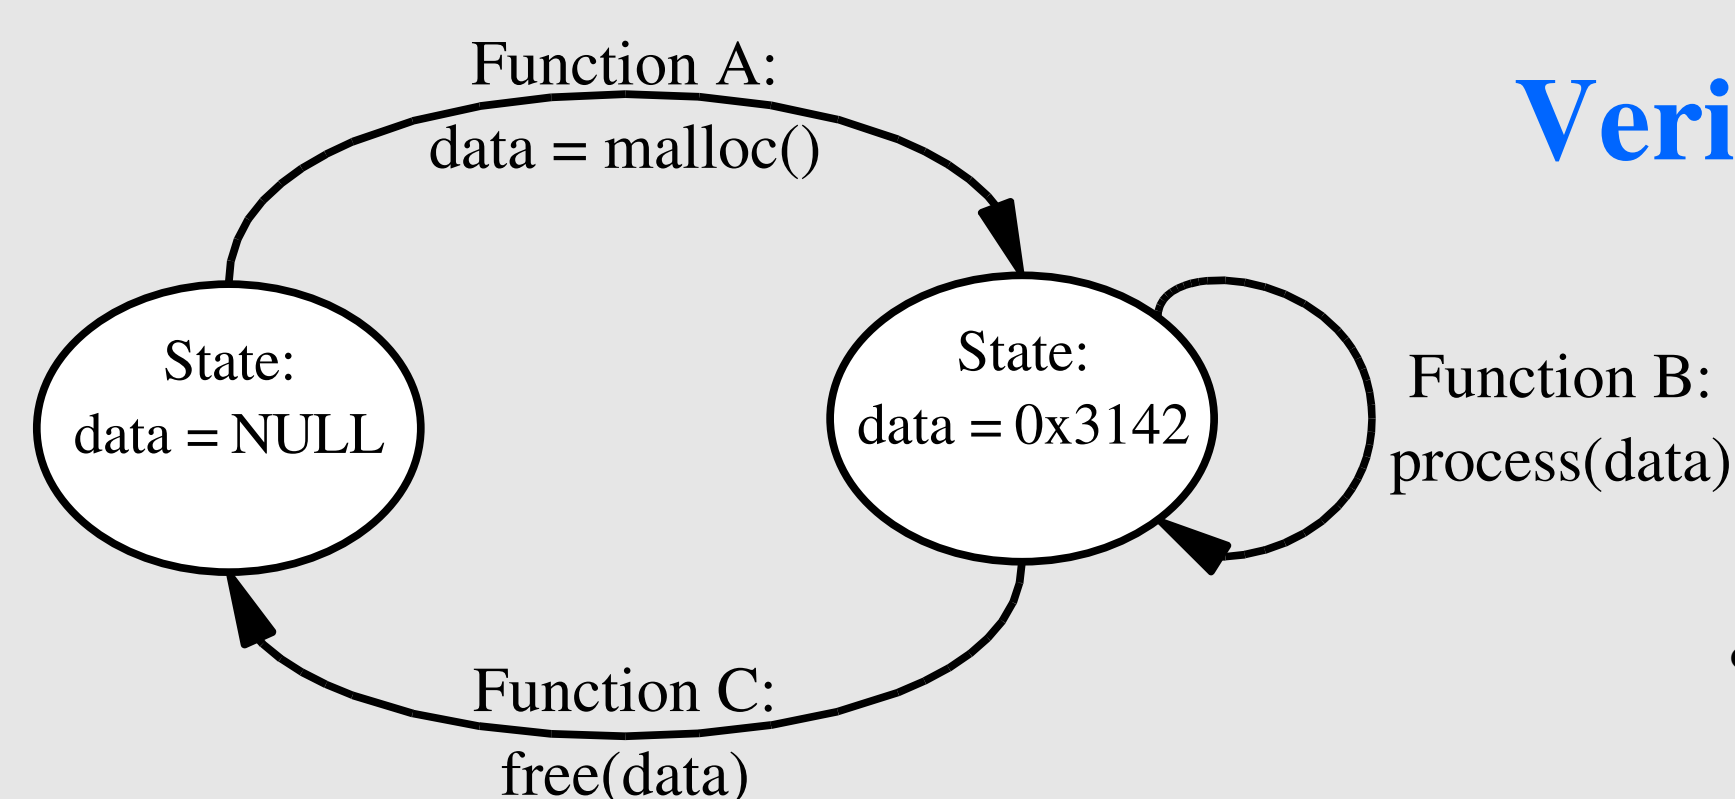
- Difficulties developing sensor network applications
  - Limited insight into a running system
  - Common failure mode is simply “not working”
  - Failures are not well contained
- Desire gaining insight into bottom tier sensing devices by applying static analysis

#### Bottom Tier Sensing Devices are Cool

- Developer can look at the entire system
- All-in-one hardware is great for hobbyist and students
- Show you amazing things about the world we live in

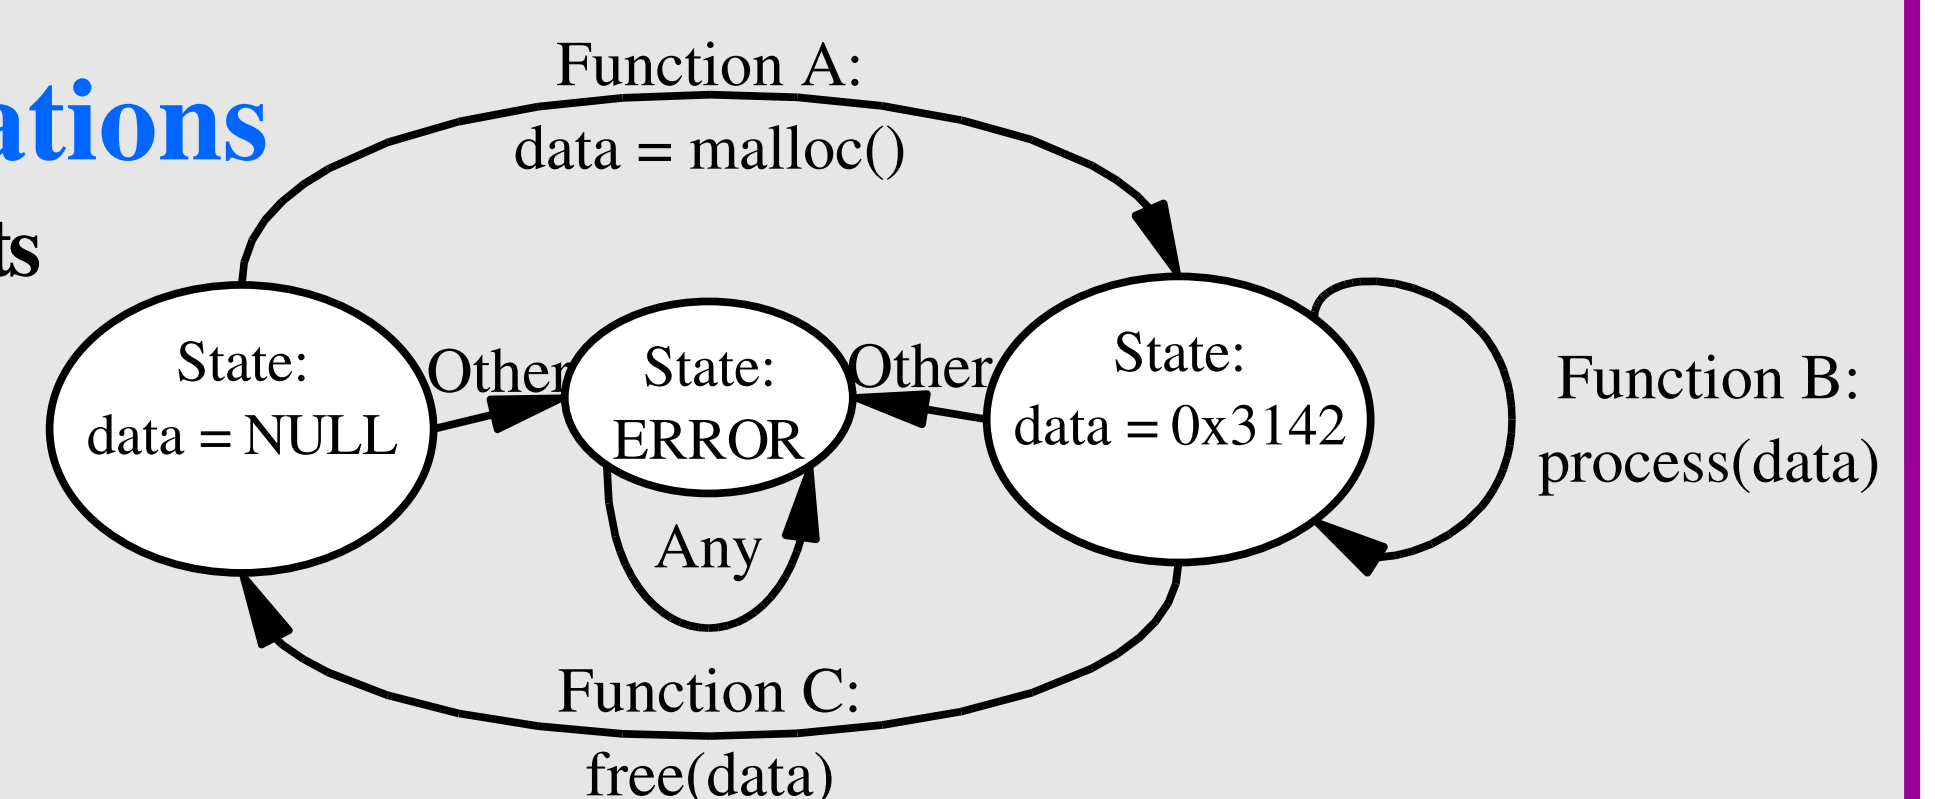


### Problem: Mapping Traditional Analysis Onto Very Reactive Systems



#### Verifying Sensor Network Applications

- Event handlers react to external events
  - Sensor detects an event of significance
  - Message is received over the radio
  - Timer injects event into the system
- Naïve model assumes any event can occur at any time



### Proposed Solution: Drive Static Analysis with Ordering Specifications

#### Lighthouse

- Simple exclusive resource ownership model
  - All heap memory is owned by exactly one module
  - Owner is responsible for releasing the memory or transferring it to another module
  - Modules may not free or transfer ownership of heap data that they do not own
- Uses standard dataflow techniques to verify each function
  - May / must pointer analysis used to resolve aliases
  - Function summaries used to handle function calls

Verified Memory Leaks	8
False Memory Leaks	8
Verified Dangling Pointers	0
False Dangling Pointers	9

Warnings generated from the analysis of 213 unique versions of SOS modules totalling 28042 SLOC

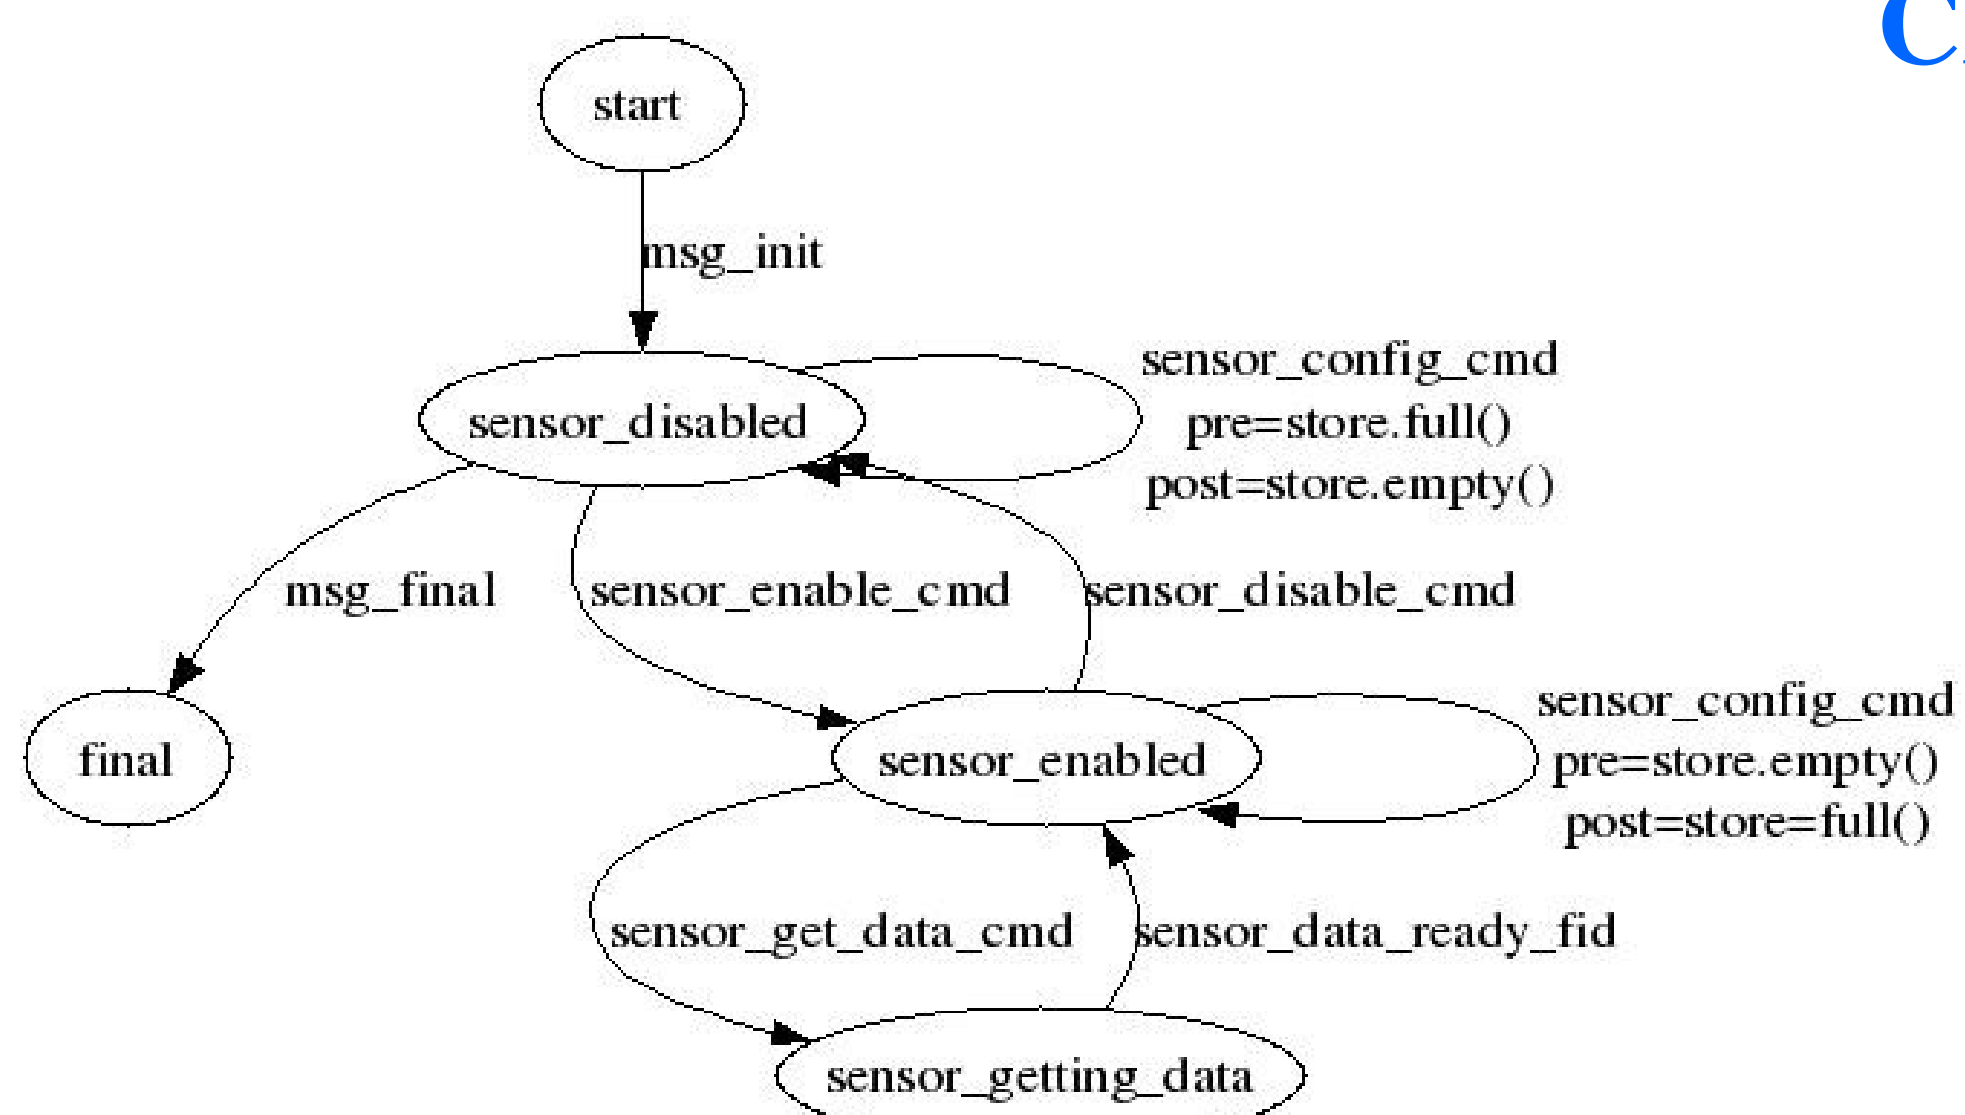
Verified Memory Leaks	2
False Memory Leaks	22
Verified Dangling Pointers	0
False Dangling Pointers	11

Warnings generated from the analysis of 40 source files making up the SOS kernel totalling 9223 SLOC

```
mod_op = (sos_module_op_t*) ker_msg_take_data(msg);
if(mod_op == NULL) return -ENOMEM;
if(mod_op->op == MODULE_OP_INSMOD) {
    existing_module = ker_get_module(mod_op->mod_id);
    if(existing_module != NULL) {
        uint8_t ver = sos_read_header_byte(...);
        if (ver < mod_op->version) {
            ker_unload_module(...);
        } else {
            return SOS_OK;
        }
    }
    ret = fetcher_request(KER_DFT_LOADER_PID,
        mod_op->mod_id, mod_op->version,
        ntohs(mod_op->size), msg->saddr);
    s->pend = mod_op;
    ker_led(LED_RED_TOGGLE);
    return SOS_OK;
}
```

Error: Expression mod\_op is not stored after instruction #line 125  
 mod\_op = (sos\_module\_op\_t\*)  
 ker\_msg\_take\_data((unsigned char)18, msg);

#### Checkpoint



- Static analysis driven by finite state machine specification
- Specification provided by the developer describes correct execution of a program
  - Set of arbitrary states associated with specific source code
  - Set of events that cause transitions between states
  - Persistent store pre- and post- conditions for stores when a specific event causes a transition between two states
- Critical for the reactive sensing environments
- Checkpoint used to refine the Lighthouse analysis

#### Future Work

- Specification inference
  - Infer memory manipulation within individual functions
  - Infer interfunction pre- / post- state dependencies
- Finite state machine specification
  - What is the best way to describe the FSM?
  - Transitioning from functional specification to slice-oriented specification